

Format & length

We're flexible on format. The cap we're really setting is on **judge attention budget**, not artifact size — submissions can include arbitrarily large code, knowledge bases, or formal specs, as long as you tell judges which slice deserves their attention first.

Reading budget: up to ~10 pages of end-to-end engagement, allocated however suits your submission — across written content, a knowledge-artifact slice, pseudocode, or a mix. Beyond this, judges skim and inspect particular areas they choose to dive into.

Written submissions (specs, protocols, analyses, critiques):

- **Core** (within the reading budget): $\leq \sim 10$ pages
- **Supporting material**: up to ~ 20 more pages.
- **Appendices**: no limit. Formal type definitions, schemas, grammars, ontologies, worked examples, and reference material can live here.

Code / executable submissions:

- No length limit on the code itself.
- Include **either** readable pseudocode (≤ 2 pages capturing algorithm and key decision points) **or** a $\sim$ one-click runnable demo (single command via Docker / Colab / Modal / similar; fresh machine to running process in ~ 5 min). Both is better.
- Short README pointing judges at the most informative invocation(s) to try and the most interesting parts of the code to inspect.

Knowledge artifacts (structured outputs your method produces):

- No 'length' limit on the artifact itself.
- Designate **representative slices** that fit within the reading budget alongside any written core, perhaps formatted as a brief tour of a few select entry points.

AI-assisted reading is welcome on both sides. Include your own AI-generated digest of long artifacts if useful; judges may also use AI to triage.