

How we'll judge. We score submissions on the dimensions below. Strong submissions don't need to land well on every dimension — a deep contribution on a few can still win the top tier — but they should plausibly engage most of them.

1. Epistemic uplift. Does this actually help a thoughtful person reason better about the case?

- Meaningfully better than off-the-shelf deep research / top-of-range Claude Code investigation on the same sub-questions
- Faithful to evidence — doesn't sand off uncertainty or smuggle in unsupported confidence
- Makes load-bearing evidence visible: which pieces are actually driving the conclusion? Which claims and evidence are getting unearned, or not enough weighting in an analysis?
- Surfaces what matters: cruxes, missing perspectives, rhetorical-vs-evidential moves

2. Generalizability. Will the workflow travel?

- Works across cases of different shape (curated debate / confident answer w/ complex evidence / mundane-but-contested)
- Plausibly applicable to cases beyond the three we provide
- Not narrowly overfit to the provided case studies

3. Compounding & shareability. Do the artefacts produced help future investigators build on this work?

- Outputs are structured / interrogable, not just narrative summaries
- Another team could pick up the artefact and extend it
- Pieces could plausibly interoperate with other submissions' pieces

4. Scalability. Does the approach get better with more compute, better models, or more contributors?

- Not bottlenecked on any single hand-designed human step
- Benefits as base-model capability rises
- Benefits from increased resources: more (adversarial) scrutiny, incremental sources, more effort/compute spent on checks

5. Methodological transparency. Is the submission well-specified enough to evaluate, replicate, and critique?

- Spec/workflow written down, with the key decisions and tradeoffs called out
- Where the creator is uncertain, that uncertainty is named — not papered over
- A judge can tell why the methodology is shaped this way, not just what it does

6. Adversarial robustness. How well do the artefacts and methodology hold up when participants and consumers have differing views and priorities?

- Outputs withstand motivated reading and downstream-model interrogation
- The methodology resists being gamed by sources optimizing to mislead
- Failure modes and uncertainties are named and bounded, not hidden

7. Insight contribution. Does the submission shift how we think about the problem itself?

- Surfaces sub-problems, framings, or considerations we'd missed
- Critiques that force re-evaluation of promising approaches, especially with potent counterexamples
- Comparative analyses that surface non-obvious tradeoffs or cutting dilemmas across methodologies

Prize tiers (guidance, not a formula — judges' discretion):

Tier	Range	Looks like
Transformative	\$35k–\$50k	Substantial advance on the SOTA — reshapes how we think the next generation of this tooling should be built. Could be a single deep contribution or a broader push across several dimensions.
Strong	\$15k–\$35k	Notable improvement on the SOTA. Either a meaningful gain across several dimensions, or a clearly impressive gain on one.
Promising	\$5k–\$15k	A real but mild improvement on the SOTA, or a partial/exploratory contribution containing insight or working components we'd want to build on.

Multiple prizes possible per tier; pool can expand for a wave of strong work; strong submissions may also lead to offers of further funded work with us.

Notes for judges.

1. Anchor against good baselines. Before scoring, check what off-the-shelf deep research or a careful Claude Code investigation produces on the same sub-question. The bar is "meaningfully better than that."
2. Read for the spec, not the polish. A clear workflow with a rough prototype usually beats a polished prototype with opaque methodology.
3. Run it, don't just read it. For tool and workflow submissions, exercise the methodology on a sub-question you're personally curious about — usefulness shows up in use, not in the writeup.